

Machine Learning Report

Raphael Barman, Hakim Invernizzi, Rehan Mulakhel
EPFL students

I. PART I: MACHINE LEARNING METHODS IMPLEMENTATION

Here is the description of our implementation of the six machine learning methods.

Please note that we defined our own metric to estimate prediction error, which is not the MSE but simply the percentage of incorrectly predicted outputs, which we call inaccuracy. This metric is used exclusively to evaluate performance and is not part of the implementation.

A. Linear regression using gradient descent

The implementation is simple: we iteratively update the initial weight by subtracting a pondered gradient. Figure 1 (a) shows the cross-validation results with 500 iterations and gamma taking value in the range $[10^{-3}, 10^0]$ with step 10. It can be noted that the method gives around 66% inaccuracy in the best case, and performs best when gamma is on the order of 10^{-1} .

B. Linear regression using stochastic gradient descent

The difference compared to the previous method is that at each iteration the dataset is accessed at random indexes in order to create size-1 batches. The batch gradient is then computed, pondered and used to update the weight. Figure 1 (b) shows the cross-validation results with 10000 iterations and gamma taking value in the range $[10^{-6}, 10^{-1}]$ with step 10. it can be noted that the method also gives 66% inaccuracy in the best case, and performs best when gamma is on the order of 10^{-3} .

C. Least squares regression

The pseudo-inverse of the feature matrix tx is computed and then multiplied with the classification vector y to obtain the weight. The inaccuracy is of around 25.5% for both train and test dataset.

D. Ridge regression

The weight is computed by solving the ridge regression equation, yielding $w = (X^T X + \lambda I)^{-1} X^T y$. Figure 1 (c) shows the cross-validation results with lambda taking value in the range $[10^{-8}, 10^0]$ with step 10. It can be noted that the performance of the ridge regression is comparable to that of the simple least square regression, suggesting that regularisation applied to the whole dataset provides little improvement. This might mean that the problem is well-conditioned, or in other words that the model isn't over- nor under-fitting.




Figure 1. Result of Cross-Validation with 10 folds. On the X-axis we have the ML method parameter to be tuned, on the Y-axis a measure of the inaccuracy. From top to bottom: (a) least squares GD (b) least squares SGD (c) ridge regression (d) logistic regression (e) logistic regression with regularization

E. Logistic regression

The weight is updated by subtracting a gradient whose ponderation takes into account its hessian matrix. Figure 1 (d) shows the cross-validation results with 10000 iterations and gamma taking value in the range $[10^{-10}, 10^{-1}]$ with step 10. It can be noted that method gives around 25.7 % inaccuracy in the best case.

F. Regularized logistic regression

Logistic regression with the addition of a penalty term to account for linearly separable data. Figure 1 (e) shows the cross-validation results with 10000 iterations, gamma on the order of $8 \cdot 10^{-3}$ and lambda taking value in the range $[10^{-10}, 10^{-1}]$ with step 10. It can be noted that the method gives around 26.25% inaccuracy in the best case, reinforcing what previously said about regularisation.

II. OUR MODEL

A. Exploratory Data Analysis and Feature Processing

1) *Null values*: The presence of null values is well defined and allowed us to separate the dataset in 6 groups. We noticed that all null values could be explained by the number of jet (PRI_jet_num) and by the presence or not of a measure for the mass (DER_mass_MMC), thus we created 6 groups, for number of jet $\{0, 1, 2 - 3\}$ and the presence or not of mass measure.

2) *Percentiles*: By looking at the 95th percentile and the maximum values of each feature, we saw that in most cases, there were certainly a lot of outliers in some features.

3) *Histograms*: We used histograms to have an idea of the underlying distribution.

Using the histograms, we saw that the features related to 'phi' (PRI_tau_phi, PRI_lep_phi, PRI_met_phi, PRI_jet_leading_phi, PRI_jet_subleading_phi) have an uniform distribution that is the same for 'signal' and for 'background', this could mean that these features do not add any relevant information for the classification.

We also noticed that some features (DER_mass_MMC, DER_mass_transverse_met_lep, DER_mass_vis, DER_pt_h, DER_mass_jet_jet, DER_pt_tot, DER_sum_pt, DER_pt_ratio_lep_tau, PRI_tau_pt, PRI_met, PRI_met_sumet, PRI_jet_leading_pt, PRI_jet_subleading_pt, PRI_jet_all_pt) have a long tail or are exponential, thus it could be a good idea to log-normalize them before using them.

B. ML method and Cross-validation

We chose to use Ridge Regression as the baseline since it has the lowest inaccuracy. All inaccuracy results that are given have been made on a 10-fold cross validation using Ridge Regression. The baseline, i.e. Ridge Regression on a ten-fold, is of inaccuracy 0.255152 with $\lambda = 10^{-10}$. We will now explore five different ways of improving the baseline.

1) *Group separation*: Using the observation we made above, we separated the dataset in 6 groups. Our hope was that since we do not have to deal with null values, the algorithms would find a better approximation of the underlying distribution. We noticed an improvement over the baseline: an inaccuracy of 0.234424 with $\lambda = 10^{-10}$.

2) *Percentile cut*: We tried to remove all values above a certain percentile and replace them with the percentile, in order to delete all outliers and allowing to have a less complex model. Using the 95th percentile, we obtained an improvement: an inaccuracy of 0.248784 with $\lambda = 10^{-10}$.

3) *Log-normalization*: Since some features had the shape of power laws and/or exponentials, we tried to log-normalize where possible (we applied log on any value > 0 in these features). However, it did not result in an improvement over the baseline, with an inaccuracy of 0.2618 with $\lambda = 10^{-10}$. This could be explained by the need of keeping the same distribution shape in order to be able to differentiate 'signal' and 'background'.

4) *Removing features*: Since we noticed that the shape of all features related to 'phi' were uniforms, we tried to remove those features in order to reduce the model complexity. We got a slight improvement, with an inaccuracy of 0.252068 and $\lambda = 10^{-10}$.

5) *Features augmentation*: In order to account for the model complexity, one can add some polynomial basis to improve the fit. We obtained an improvement with a polynomial basis of degree 3: inaccuracy = 0.229108, $\lambda = 10^{-10}$. The improvement is quite important, this shows that the underlying model is certainly of higher degree than one.

6) *Final Model*: We tried different combination of the 4 improvements we found above and found the best result with a mix of group separation, percentile cut and polynomial basis. We found for results: inaccuracy of 0.1681 with parameters for each groups: $(d_0 = 7, \lambda_0 = 0)$ $(d_1 = 5, \lambda_1 = 0)$ $(d_2 = 9, \lambda_2 = 10^{-4})$ $(d_3 = 4, \lambda_3 = 1.66 \cdot 10^{-8})$ $(d_4 = 8, \lambda_4 = 4.64 \cdot 10^{-4})$ $(d_5 = 4, \lambda_5 = 0)$.

Note that until now, the usage over the Ridge Regression over least squares did not make much sense since we always had negligible lambdas, however, with our final model, some group benefits from the ridge regression.